Spectral Filtering for Vision–Language–Action Models

Shivam Kak [sk3686@princeton.edu] [1] Ishan Saha [is1893@princeton.edu] [1] Connor Brown [cb6856@princeton.edu] [1]

[1]Princeton University

PRINCETON UNIVERSITY

Abstract

We integrate Spectral Transfer Units (STUs) [2] into the $\pi_{0.5}$ flow-matching Vision–Language–Action model [1] and evaluate them on LIBERO [5] action prediction across *four* architectural placements: post-hoc on the action expert; pre-input on the noisy action chunk; on the predicted velocity; and pre-input on the PaliGemma prefix. We compare against a parameter-matched diagonal-SSM (Mamba) control to isolate Hankel-specific structure. Across all placements, all spectral-filter counts K, all action horizons $H \in \{5, 10, 20, 50\}$, and both warmup schedules, STU ties or underperforms the baseline. Mamba ties STU exactly, indicating the Hankel eigenbasis is not buying anything beyond generic linear recurrence in this regime. A short-warmup probe rules out the budget-bound hypothesis: forcing the spectral block to grow degrades the model. The result is a comprehensive negative finding that delineates where Hankel-eigenbasis filtering does and does not transfer.

Background: $\pi_{0.5}$ + Spectral Filtering

$\pi_{0.5}$ flow-matching VLA. A SigLIP vision encoder + PaliGemma language model + a small action expert. An action chunk $a \in \mathbb{R}^{H \times D_a}$ is produced by integrating a learned velocity field

$$x_{t+\mathrm{d}t} = x_t + \mathrm{d}t \cdot v_\theta(x_t, t), \quad x_0 \sim \mathcal{N}(0, I).$$

The action expert sees the full chunk in parallel, so causal attention is the only mechanism mixing information across the H time positions of the chunk—a natural fit for spectral methods.

Spectral Transfer Units (STU). Convolve along the horizon with eigenvectors ϕ_k of the Hankel matrix $Z_{ij} = \frac{2}{(i+j)^3-(i+j)}$, scaled by the fourth root of their eigenvalues:

$$y = \sum_{k=1}^{K} \Big((u * \phi_k^+)\, M_k^+ \;+\; (u * \phi_k^-)\, M_k^- \Big)$$

where $\phi_k^\pm$ split filters into even/odd-symmetric components, $M_k^\pm$ are learnable, FFT in $O(L \log L)$.

Mamba-style control. Parameter-matched diagonal SSM $\dot{x} = Ax + Bu$, $y = Cx$ with HiPPO-LegS init. Isolates whether *Hankel-specific* structure helps vs. generic linear recurrence (per midterm reviewer ask).

Architectural Placements

v1 (post-hoc, action expert): residual STU *after* the action-expert transformer, randomly initialized.
v2 (pre-input, action expert): STU on the noisy action chunk x_t *before* the transformer; zero-init out-projection so step 0 is bit-identical to baseline.
v3 (output velocity): STU on predicted $v_t \in \mathbb{R}^{B \times H \times 32}$, the natural axis for a "smoothness" prior.
v4 (PaliGemma prefix): mirror of v2 on the VLM side; STU on the prefix embeddings (∼968 image+language tokens) before the VLM forward, zero-init out-projection.

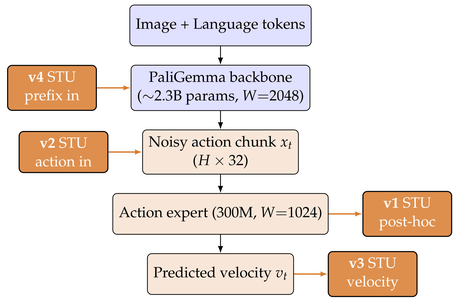

Figure 1. The four sites in $\pi_{0.5}$ where we tested an STU residual block. v1 is the original midterm design; v2/v3/v4 share the dreamer-aligned recipe (pre-input or pre-output, zero-init out-projection). Each placement isolates a different question: whether residual STUs help *after* a converged action expert (v1), at the chunk-input or velocity-output stages where horizon mixing is most plausible (v2/v3), or on the much larger vision-language backbone where capacity may be less of a bottleneck (v4).

Experimental Setup

Training. Fine-tune from $\pi_{0.5}$ *base* (`gs://openpi-assets/checkpoints/pi05_base`)—*not* the LIBERO-finetuned checkpoint. AdamW 5×10^{-5}, cosine decay (10k or 1k warmup), batch 8, gradient clip 1.0, bfloat16, 5000 steps per run.

Compute. 26 LIBERO fine-tunes, ∼30 GPU-hours on Princeton's `della` cluster (H100 NVL 95 GB and 2× RTX PRO 6000 Blackwell 97 GB).

Evaluation. Last-100-step mean training loss (smoothed flow-matching MSE) is the reported metric; differences below ±1% are within run-to-run noise.

Results: H=10 Architecture Comparison

Model	Final Loss	Δ Base	Params
Baseline	**0.0299**	—	0
STU-K4 (post-hoc)	0.0316	+5.7%	8.4M
STU-K8 (post-hoc)	0.0313	+4.7%	16.8M
Mamba-K4 (post-hoc)	0.0311	+4.0%	8.4M
Mamba-K8 (post-hoc)	0.0312	+4.3%	16.8M
STU-v2-K4 (act-input)	**0.0298**	**−0.5%**	8.4M
STU-v2-K8 (act-input)	**0.0298**	**−0.4%**	16.8M
STU-v3-K4 (output vel)	**0.0297**	**−0.6%**	0.07M
STU-v4-K4 (PaliGemma)	**0.0296**	**−1.2%**	33.6M
STU-v4-K8 (PaliGemma)	**0.0296**	**−1.0%**	67.1M

Table 1. Last-100-step mean loss at H=10. Post-hoc residuals *hurt*; all three zero-init pre-input placements (v2, v3, v4) recover baseline within run-to-run noise.

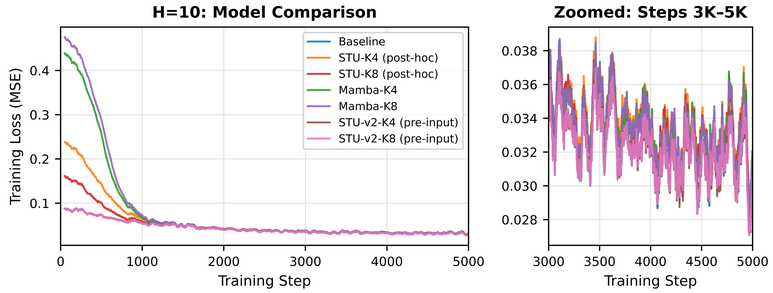

Figure 2. H=10 training loss (50-step rolling mean). Post-hoc variants (orange/green) start with elevated loss and only catch up by step ∼2000. Zero-init STU-v2 (brown) overlaps baseline (blue) from step 0—the visual signature of zero-init.

Horizon Ablation

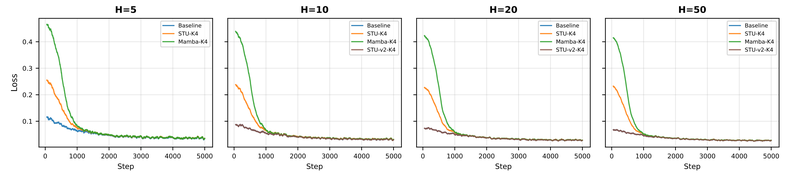

Figure 3. Horizon sweep ($H \in \{5, 10, 20, 50\}$). Post-hoc variants elevated at every horizon. STU-v2 (brown) overlaps baseline from step 0 at every horizon. Longer H reaches lower terminal MSE because each loss value averages more correlated time-steps.

H	Base	STU-K4 post-hoc	STU-v2-K4	STU-v4-K4
5	0.0338	0.0352 (+4.2%)	—	—
10	0.0299	0.0316 (+5.7%)	**0.0298**	**0.0296**
20	0.0276	0.0286 (+3.5%)	**0.0275**	**0.0275**
50	0.0269	0.0274 (+1.9%)	**0.0270**	—

Table 2. Horizon ablation, all warmup-10k. Post-hoc gap shrinks (+5.7% @ 10 → +1.9% @ 50) but does not close. Linear extrapolation crosses zero only at $H \approx 80$–100, beyond LIBERO. Zero-init v2 & v4 tie baseline at every horizon tested.

Spectral Weight Analysis

We loaded the v2 checkpoints and inspected $M_k^\pm$ directly to verify the spectral block is being trained:

$$\max |M_k^\pm| \approx 8\times10^{-3}, \quad \overline{|M_k^\pm|} \approx 7\times10^{-4}$$

Frobenius norms 0.29–0.49 across the full $K \times D_a \times W$ tensor. Default random-init scale would be $(KD_a)^{-1/2} \approx 0.088$, so after 5k steps $M_k^\pm$ has grown to about **1% of a typical layer's magnitude**. The optimizer has measurably moved $M_k^\pm$ off zero, but not enough to materially alter the network's output.

Warmup Schedule Sensitivity

Shortening LR warmup from 10k to 1k steps causes $\|M^+\|_F$ to grow from 0.36 to 0.92 (2.5×), forcing the spectral block off zero-init faster.

	Base	v2-K4	v3-K4	v4-K4
warmup-10k	0.0299	−0.5%	−0.6%	−1.2%
warmup-1k	0.0277	+1.5%	+2.0%	+1.9%

Table 3. Warmup probe at H=10, all three zero-init placements. Absolute loss improves with shorter warmup at every site, but the relative position of every STU variant *worsens* when $M_k^\pm$ is allowed to grow—consistent across all three placements. **This rules out the budget-bound explanation:** more learning signal → *worse*, not better.

Key Findings

1. **Negative result is robust across 4 placements** (post-hoc, action-input, output-velocity, VLM-prefix), all K, all H, both warmups. Spectral filtering does not add a useful inductive bias for $\pi_{0.5}$ action prediction on LIBERO.
2. **Hankel structure is irrelevant here.** Mamba ties STU at every configuration tested, so the Hankel eigenbasis is not buying anything beyond generic linear recurrence.
3. **The horizon hypothesis is directionally right but quantitatively weak.** The post-hoc gap shrinks with H, but linear extrapolation crosses zero only at $H \approx 80$–100, beyond LIBERO.
4. **Placement and zero-init matter more than the spectral basis.** The pre-input zero-init recipe rescues the negative result everywhere (no longer hurts), but never produces a positive one.
5. **Even on the much larger PaliGemma backbone (v4),** with 25× more parameters than the action expert, the result is unchanged. The headroom hypothesis is not the explanation.

Discussion

Why doesn't the spectral block engage? The dreamer-aligned recipe makes the residual branch start at zero, but it does not give it any reason to leave zero. With H=10 noisy chunks predicted from a near-converged backbone, the per-step MSE gradient already lives in the action expert's existing weight space; the optimizer has no incentive to route it through a freshly-initialised parallel pathway. The 1k-warmup probe makes this concrete: when we force $M_k^\pm$ to grow 2.5×, performance *worsens* at every placement.

The horizon hypothesis is real but small. The post-hoc gap shrinks monotonically from +5.7% at H=10 to +1.9% at H=50, validating the directional claim from our midterm. But linear extrapolation crosses zero only at $H \approx 80$–100, well beyond what LIBERO action chunks support. The spectral inductive bias has the right shape; the wrong magnitude.

Architecture matters, structure does not. Mamba-K{4,8} ties STU at every configuration, so the cost—and the benefit—of inserting a new linear-recurrence branch is dominated by *where* you put it and *how* you initialise it, not by the choice of Hankel eigenbasis vs. HiPPO-LegS diagonal SSM.

Future Work

- Success rates on the 130 LIBERO tasks (we measure only training MSE; rollout behavior could differ even at near-equal MSE).
- A generic-SSM control on the VLM side (Mamba-v4): the FFT-based diagonal-SSM kernel exceeds 32 GB on the 968-position prefix at W=2048; a chunked-FFT or selective-scan re-implementation would close this gap.
- STU-v3 at $H \in \{20, 50\}$ and STU-v4 at H=50 (compute-bound).
- A learned residual gate, to verify the optimizer is not actively zeroing the block.
- Pre-training a $\pi_{0.5}$ checkpoint *from scratch* with the spectral block in place, to test whether STU helps when not introduced as a residual on a converged backbone.

References

[1] Physical Intelligence. π_0: A Vision-Language-Action Flow Model. *arXiv:2410.24164*, 2024.
[2] Agarwal, N., Hazan, E., et al. Spectral State Space Models. *arXiv:2312.06837*, 2023.
[3] Gu, A., Dao, T., Ermon, S., Rudra, A., Ré, C. HiPPO: Recurrent Memory with Optimal Polynomial Projections. *NeurIPS*, 2020.
[4] Gu, A., Dao, T. Mamba: Linear-Time Sequence Modeling with Selective State Spaces. *arXiv:2312.00752*, 2023.
[5] Liu, B., et al. LIBERO: Benchmarking Knowledge Transfer for Lifelong Robot Learning. *NeurIPS*, 2023.
[6] Physical Intelligence. OpenPI. `https://github.com/Physical-Intelligence/openpi`, 2024.

Acknowledgements

We thank Professor Dhruv Shah, instructor of ECE 534, and teaching assistants Lihan Zha and Asher Hancock for their feedback that informed the direction of this project. Compute provided by the Princeton `della` cluster.